

TCP Congestion Control Comparison

Arne Esterhuizen

20 Janaury 2012

Abstract

The purpose of this report is to investigate the effects that different TCP variants have on each other. These TCP variants differ by the congestion control algorithms they employ. Congestion control algorithms affect how much network traffic is generated by TCP at any one time and aims to prevent a TCP connection from over utilising the network. This report investigates the different congestion control algorithms that are included as loadable modules in the Linux kernel, and discusses experiments performed on how these congestion control algorithms compete for network resources. It will be shown that some TCP variants are able to co-exist, whilst others tend to use excessive bandwidth, potentially smothering other competing TCP connections. Should the results prove interesting, the project may be given to third year computer network students to perform as an assignment.

1 Table of Contents

2 Introduction

The *Transmission Control Protocol* (TCP) controls how much data it transmits over a network by utilising a congestion window.

TCP cannot send more data than the congestion window allows. The size of TCP's congestion window depends on the condition of the network. In the case of the network experiencing heavy data traffic, the congestion window would be shortened. In the case of the network being only lightly loaded with data traffic, the congestion window would be expanded. How and when to adjust the congestion window depends on whatever form of congestion control a TCP connection utilises. Congestion control algorithms rely on various indicators to passively ascertain the state of the network. Packet loss is an indication that the network is being over utilised and that routers along the way are dropping IP packets due to limited buffer space. Also, routers can set special flags in an IP packet's header that will inform the receiving host that congestion is about to occur. The receiving host can then inform the the sending host to reduce its sending rate. Other methods include closely observing packet round trip times as well as determining the queueing delay, the time packets spend waiting in a router's buffer. Using one or more of these methods, congestion control algorithms are able to adjust their behaviour according to whether congestion has occurred or could occur. As will be seen, some congestion control mechanisms allow for unfair usage of network bandwidth, while others

are able to share bandwidth equally.

The different congestion control mechanisms available for use by the Linux kernel are:

- HTCP
- Hybla
- Illinois
- LP
- Vegas
- Reno
- BIC
- Westwood
- YeAH
- Cubic
- Highspeed
- Scalable
- Veno

The linux socket interface allows one to change the type of congestion control a TCP connection uses by setting the appropriate socket option.

A short description of the congestion control algorithms is given in section 3. Experiments were performed to see how these algorithms compete against each other for bandwidth. A description of these experiments is given in section 4, followed by a discussion of the results in section 5. Technical difficulties that may affect the practicality of this project as an assignment for students is described in section 6. Section 8 summarises the presented results.

3 Overview of TCP variants

4 Experiments performed

Experiments were carried out using client and server programs. A host running a client program generates data which is sent over the network to a host running the server program. A single server host receives data from various client hosts. Each client host is set to use a different congestion control algorithm. The amount of data received from client hosts is measured in megabits per second and is presented as a graph. This shows how much bandwidth the client host is able to utilise. The mean of a client's bandwidth measurements is used to determine its bandwidth usage in general. In these experiments, the highest possible throughput on the network is slightly over 100 Mbps.

The experiments were run on the Stellenbosch University network. All hosts involved were located on the same network segment. All experiments were carried out in the same manner, and all TCP variants were tested in pairs. Host A (a client) transmits data to host B (the server), after which host C (another client) also commences sending data to host B. Thus the two client hosts do not commence with data transmission simultaneously, there is always a short delay between one client beginning its transmission and the next client beginning its transmission. Therefore, the first client to commence transmission always utilises the network freely, and is only forced to share bandwidth as soon as an additional client commences transmission. This was found to have an impact on the behaviour of certain TCP variants, as their bandwidth utilisation appeared to change depending on

the order in which they were allowed to start (refer to section 5.3).

Every TCP congestion control algorithm was tested against all other algorithms, including itself. This was done in order to establish whether TCP variants could be placed into groups in which all members are well behaved.

5 Results

As mentioned in section 4, all TCP variants were compared to each other. The results of these experiments will be discussed in this section for each TCP variant individually. To avoid filling up this report with over 120 images, only the most interesting results will be illustrated by showing the graph produced by the experiment.

5.1 HTCP

TCP HTCP is able to co-exist with itself, with each HTCP client having a mean throughput of 52 Mbps. HTCP has the lower throughput against Illinois, BIC, Westwood, and Scalable. Against Illinois, BIC, and Scalable, HTCP has a mean throughput in the range of 38 to 41 Mbps. Against Westwood, HTCP has a throughput of 19 Mbps. HTCP is stronger against Hybla, LP, Vegas, Reno, YeAH, Cubic, and Highspeed. Against these it maintains a mean throughput in the range of 64 to 67 Mbps. The exception is vegas, against which HTCP has a mean throughput of about 95 Mbps.

5.2 Hybla

Referring to table 2, Hybla seems to be able to co-exist with itself, LP, Reno, Cubic, and

Mean Throughput (Mbps)		
HTCP	Other	
52	52	HTCP
66	37	Hybla
38	65	Illinois
65	38	LP
95	10	Vegas
66	37	Reno
41	62	BIC
19	86	Westwood
64	40	YeAH
67	37	Cubic
64	40	Highspeed
41	64	Scalable

Table 1: Mean throughput of HTCP against other TCP variants.

YeAH. Hybla seems to do badly against Illinois, BIC, Westwood, and Scalable.

Mean Throughput (Mbps)		
Hybla	Other	
37	66	HTCP
52	51	Hybla
33	70	Illinois
48	54	LP
94	10	Vegas
49	54	Reno
28	75	BIC
14	90	Westwood
46	57	YeAH
53	50	Cubic
40	63	Highspeed
29	74	Scalable

Table 2: Mean throughput of Hybla against other TCP variants.

5.3 Illinois

As can be seen from table 3, Illinois is a very unfriendly congestion control algorithm, and is unable to fairly share bandwidth with any other congestion control algorithm. The only congestion control algorithm to maintain a higher mean throughput than Illinois is Westwood. Illinois has the curious trait of not being able to co-exist with even itself, something all the other congestion control algorithms are capable of. In section 4 it was mentioned that in some cases the order in which clients are allowed to begin transmission influences the bandwidth utilised. Illinois is one algorithm for which this applies. In table 3, Illinois is allowed to begin transmission first, and even after a second client begins transmitting data, Illinois still maintains a very high mean throughput. However if the order of transmission is reversed, the competing client is allowed much more bandwidth, even if it is still less than Illinois. This is illustrated by figures 1 and 2. Illinois shows similar behaviour with all other TCP variants, with the exception of Westwood and Vegas, whose mean throughput remains the same regardless of the order in which they start transmitting.

5.4 Vegas

TCP Vegas is able to co-exist with itself. Each Vegas client has a mean throughput of about 50 Mbps. Vegas performs very poorly against all other TCP variants, achieving a mean throughput of anything between 8 and 13 Mbps, depending on which TCP variant it competes against. This is demonstrated by figure ??.




Figure 1: Illinois commences transmission first, and is joined by YeAH after 10 seconds. Illinois maintains a mean throughput of 93 Mbps, while YeAH maintains a very low mean throughput of 10 Mbps.




Figure 2: YeAH is now allowed to begin transmission before Illinois. Although still less than Illinois, YeAH now maintains a much higher than before mean throughput of 37 Mbps, while Illinois having a lower than before mean throughput of 67 Mbps.

Mean Throughput (Mbps)		
Illinois	Other	
65	38	HTCP
70	33	Hybla
13	89	Illinois
95	7	LP
97	6	Vegas
95	8	Reno
93	10	BIC
34	69	Westwood
93	10	YeAH
94	9	Cubic
94	9	Highspeed
88	22	Scalable

Table 3: Mean throughput of Illinois against other TCP variants.

5.5 Scalable

6 Technical Difficulties

7 Future Improvements

8 Conclusion




Figure 3: Vegas begins transmission, followed by BIC ten seconds later. BIC immediately achieves a mean throughput of 91 Mbps, while Vegas is reduced to a mean throughput of 13 Mbps.